

Building a Distributed Data Ingestion System with RabbitMQ

Alvaro Videla - RabbitMQ



Конференция разработчиков
высоконагруженных систем





Alvaro Videla

- Works at RabbitMQ
- Co-Author of RabbitMQ in Action
- Creator of the RabbitMQ Simulator
- Blogs about RabbitMQ Internals: <http://videlalvaro.github.io/internals.html>
- @old_sound — alvaro@rabbitmq.com — github.com/videlalvaro




About Me

Co-authored

RabbitMQ in Action

<http://bit.ly/rabbitmq>



About this Talk

- Exploratory Talk
- A ‘what could be done’ talk instead of ‘this is how you do it’



Agenda

- Intro to RabbitMQ
- The Problem
- Solution Proposal
- Improvements





<https://twitter.com/spacemanaki/status/514590885523505153>



What is RabbitMQ



Multi Protocol



<http://bit.ly/rmq-protocols>



Community Plugins

<http://www.rabbitmq.com/community-plugins.html>



Polyglot



Polyglot

- PHP



Polyglot

- PHP
- node.js



Polyglot

- PHP
- node.js
- Erlang



Polyglot

- PHP
- node.js
- Erlang
- Java



Polyglot

- PHP
- node.js
- Erlang
- Java
- Ruby



Polyglot

- PHP
- node.js
- Erlang
- Java
- Ruby
- .Net

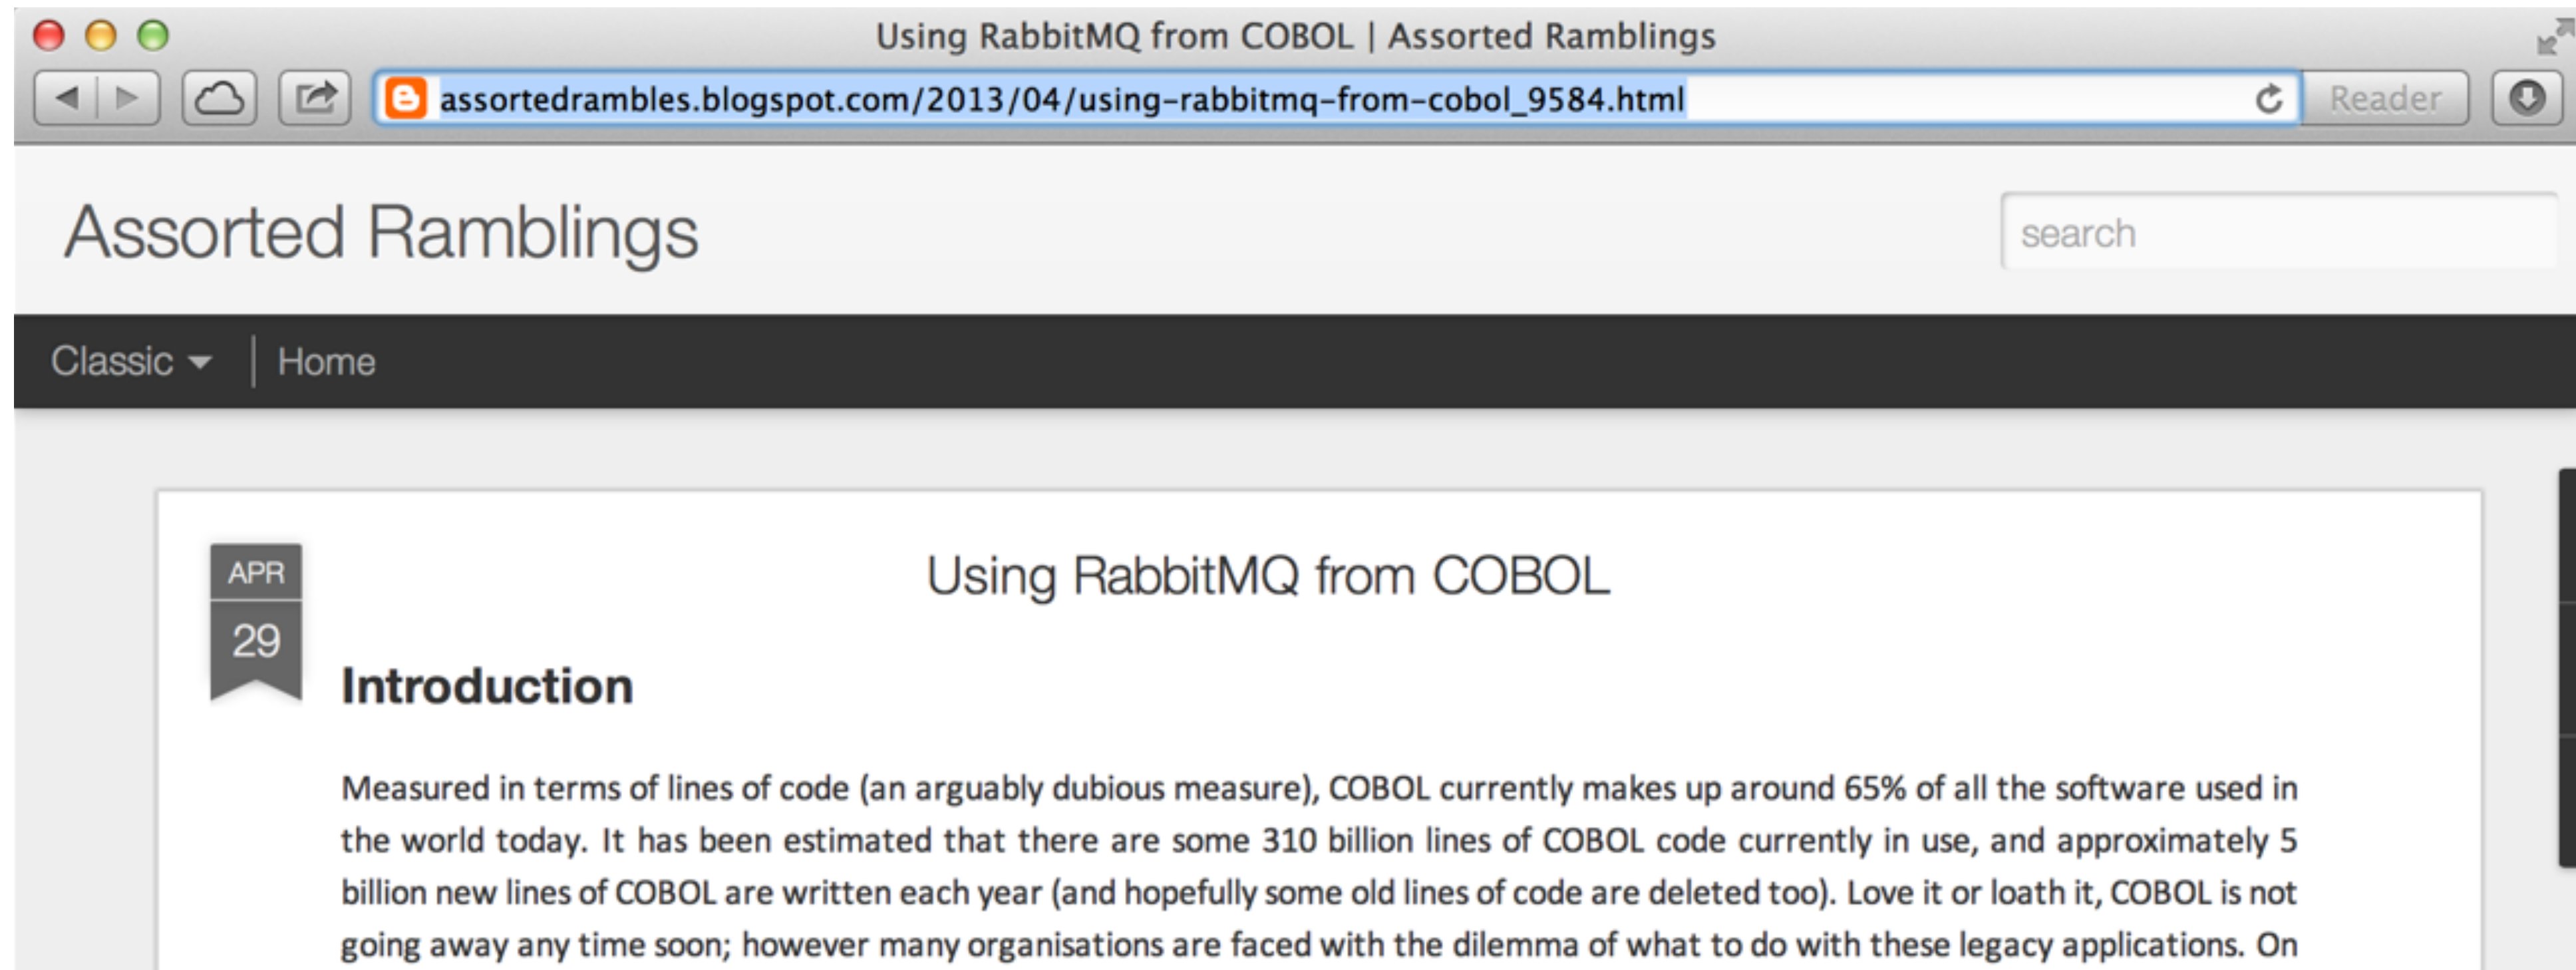


Polyglot

- PHP
- node.js
- Erlang
- Java
- Ruby
- .Net
- Haskell



Polyglot



Even COBOL!!!11



Some users of RabbitMQ



Some users of RabbitMQ

- Instagram



Some users of RabbitMQ

- Instagram
- Indeed.com



Some users of RabbitMQ

- Instagram
- Indeed.com
- Telefonica



Some users of RabbitMQ

- Instagram
- Indeed.com
- Telefonica
- Mercado Libre



Some users of RabbitMQ

- Instagram
- Indeed.com
- Telefonica
- Mercado Libre
- Mozilla



The New York Times on RabbitMQ

This architecture - Fabrik - has dozens of RabbitMQ instances spread across 6 AWS zones in Oregon and Dublin.

Upon launch today, the system autoscaled to ~500,000 users. Connection times remained flat at ~200ms.

<http://lists.rabbitmq.com/pipermail/rabbitmq-discuss/2014-January/032943.html>



<http://www.rabbitmq.com/download.html>

Unix - Mac - Windows



Messaging with RabbitMQ

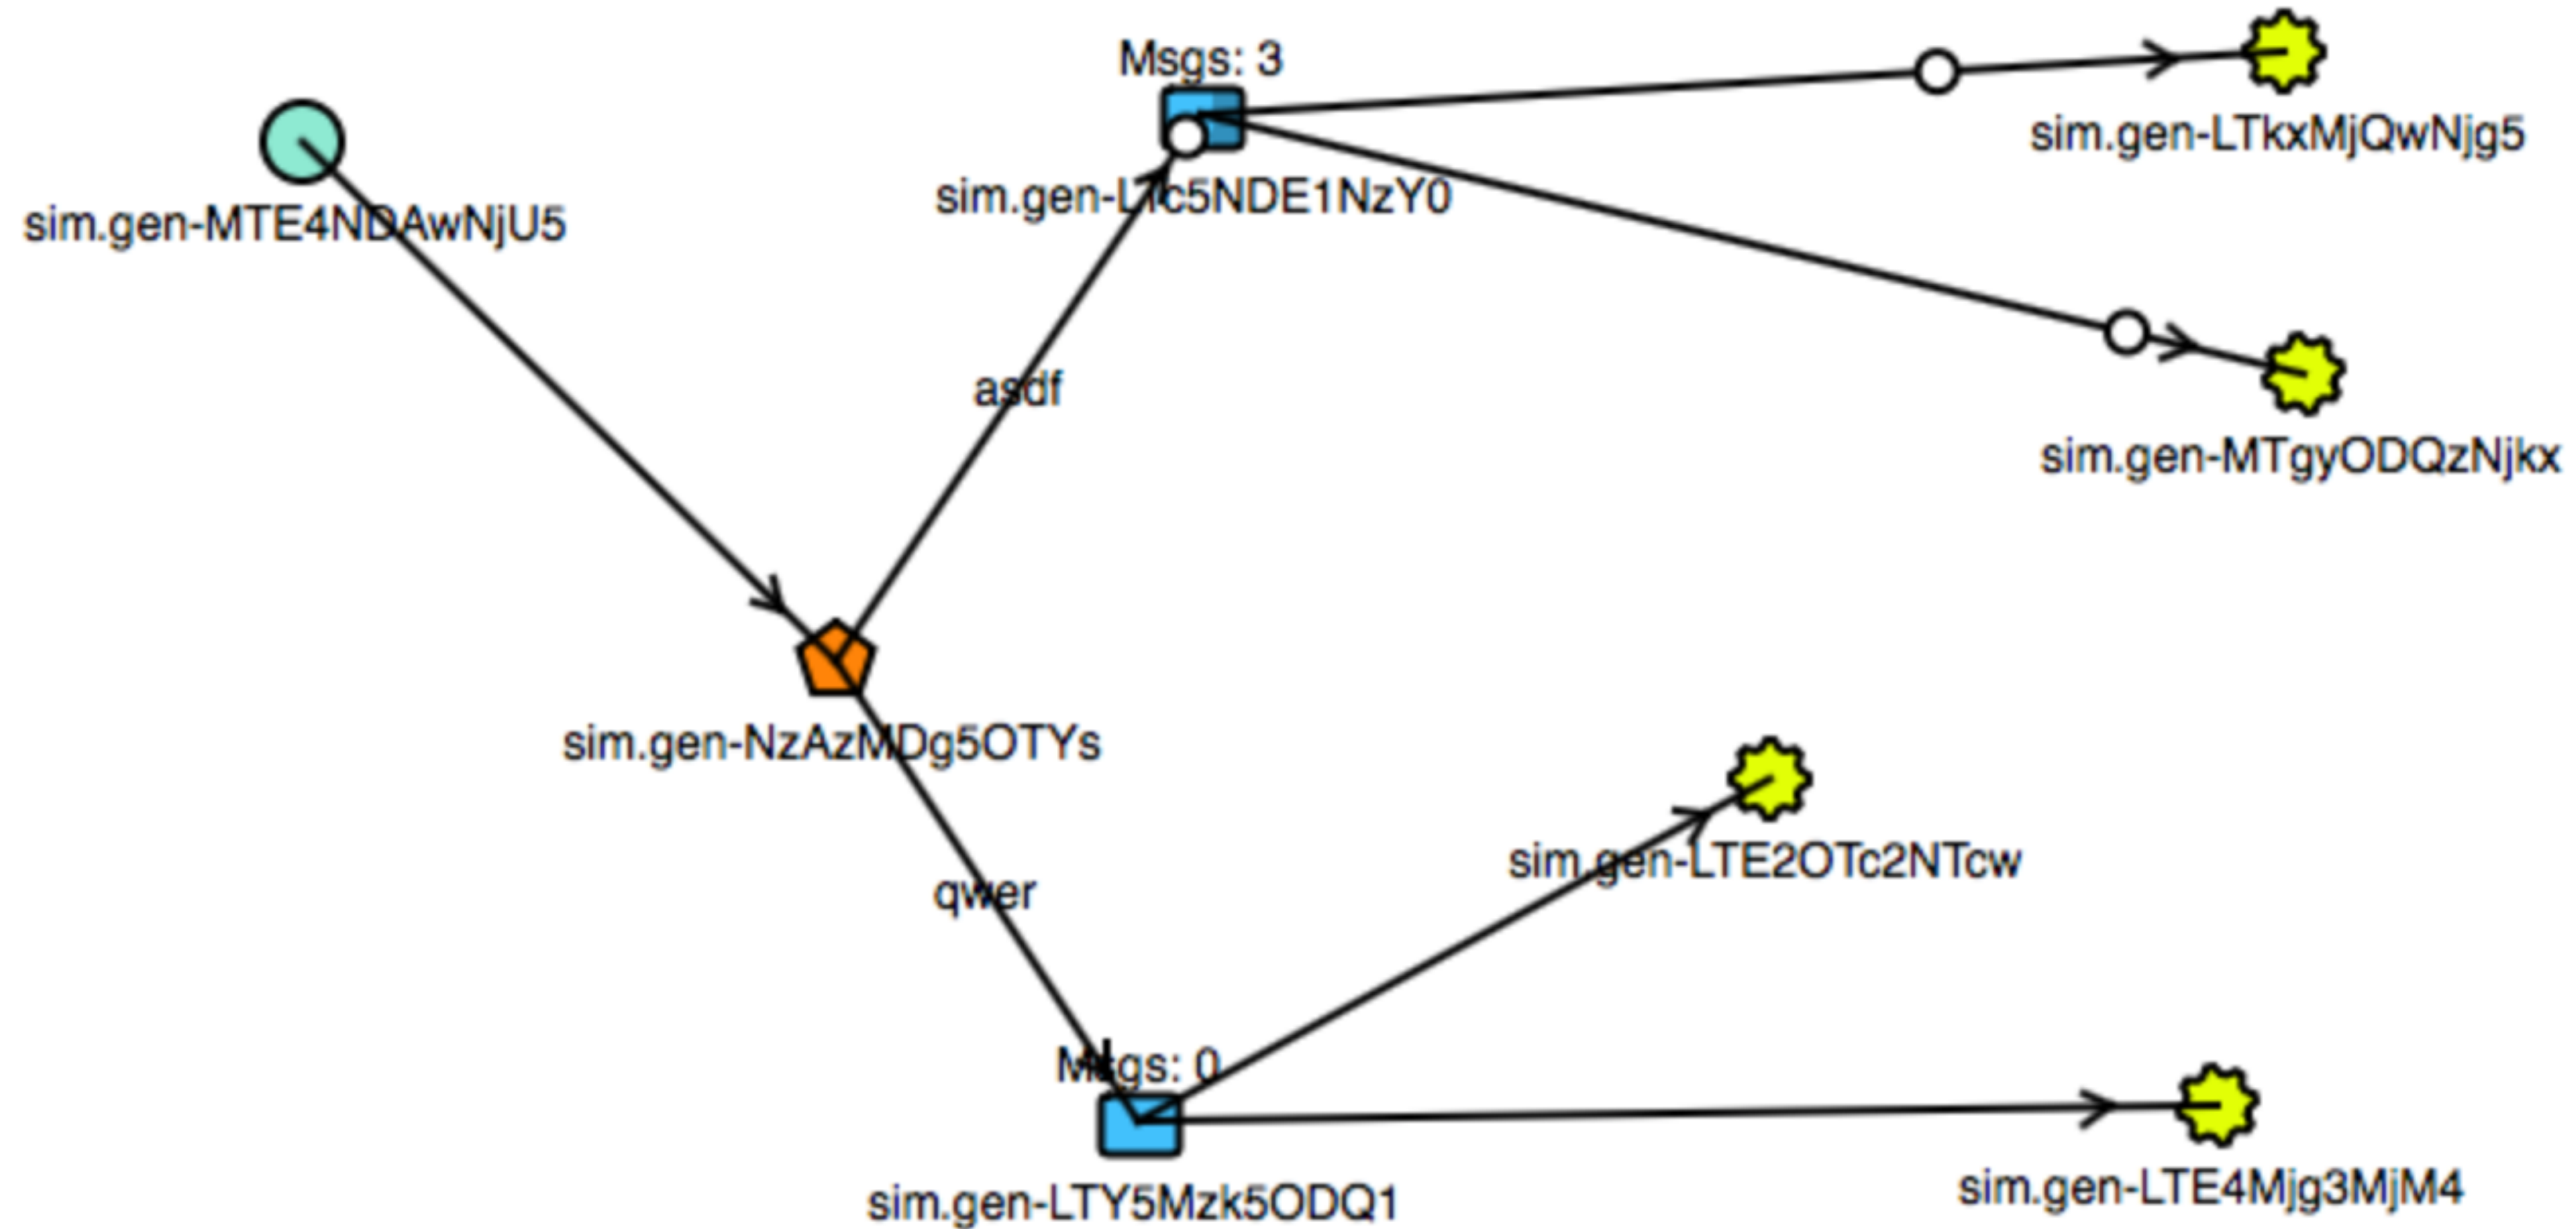
[https://github.com/RabbitMQSimulator/
RabbitMQSimulator](https://github.com/RabbitMQSimulator/RabbitMQSimulator)



<http://tryrabbitmq.com>



RabbitMQ Simulator



The Problem



Distributed Application



Data Producer

Obtain a Channel

```
{ok, Connection} =  
    amqp_connection:start(#amqp_params_network{host = "localhost"}),  
  
{ok, Channel} = amqp_connection:open_channel(Connection),
```



Data Producer

Declare an Exchange

```
amqp_channel:call(Channel, #'exchange.declare'{exchange = <<"events">>,
                                                type = <<"direct">>}),
```



Data Producer

Publish a message

```
amqp_channel:cast(Channel,  
  #'basic.publish'{  
    exchange = <<"events">>},  
  #amqp_msg{props = #'P_basic'{delivery_mode = 2},  
    payload = <<"Hello Federation">>}),
```



Data Consumer

Obtain a Channel

```
{ok, Connection} =  
    amqp_connection:start(#amqp_params_network{host = "localhost"}),  
  
{ok, Channel} = amqp_connection:open_channel(Connection),
```



Data Consumer

Declare Queue and bind it

```
amqp_channel:call(Channel, #'exchange.declare'{exchange = <<"events">>,
                                                type = <<"direct">>}),

#'queue.declare_ok'{queue = Queue} =
    amqp_channel:call(Channel, #'queue.declare'{exclusive = true}),

amqp_channel:call(Channel, #'queue.bind'{exchange = <<"events">>,
                                         queue = Queue}),
```



Data Consumer

Start a consumer

```
amqp_channel:subscribe(Channel, #'basic.consume'{queue = Queue,  
                                                    no_ack = true}, self()),  
  
receive  
    #'basic.consume_ok'{} -> ok  
end,  
  
loop(Channel) .
```



Data Consumer

Process messages

```
loop(Channel) ->  
  receive  
    {#'basic.deliver'{}, #amqp_msg{payload = Body}} ->  
      io:format(" [x] ~p~n", [Body]),  
      loop(Channel)  
end.
```



Distributed Application



Ad-hoc solution



**A process that replicates data
to the remote server**



Possible issues



Possible issues

- Remote server is offline



Possible issues

- Remote server is offline
- Prevent unbounded local buffers



Possible issues

- Remote server is offline
- Prevent unbounded local buffers
- Prevent message loss



Possible issues

- Remote server is offline
- Prevent unbounded local buffers
- Prevent message loss
- Prevent unnecessary message replication



Possible issues

- Remote server is offline
 - Prevent unbounded local buffers
 - Prevent message loss
- Prevent unnecessary message replication
 - No need for those messages on remote server



Possible issues


- Remote server is offline
 - Prevent unbounded local buffers
 - Prevent message loss
- Prevent unnecessary message replication
 - No need for those messages on remote server
 - Messages that became stale



Can we do better?



RabbitMQ Federation



RabbitMQ Federation



RabbitMQ Federation

- Supports replication across different administrative domains



RabbitMQ Federation

- Supports replication across different administrative domains
- Supports mix of Erlang and RabbitMQ versions



RabbitMQ Federation

- Supports replication across different administrative domains
- Supports mix of Erlang and RabbitMQ versions
- Supports Network Partitions




RabbitMQ Federation


- Supports replication across different administrative domains
- Supports mix of Erlang and RabbitMQ versions
- Supports Network Partitions
- Specificity - not everything has to be federated



RabbitMQ Federation



RabbitMQ Federation



RabbitMQ Federation



RabbitMQ Federation

- It's a RabbitMQ **Plugin**



RabbitMQ Federation

- It's a RabbitMQ **Plugin**
- Internally uses **Queues** and **Exchanges Decorators**



RabbitMQ Federation

- It's a RabbitMQ **Plugin**
- Internally uses **Queues** and **Exchanges Decorators**
- Managed using **Parameters** and **Policies**



Enabling the Plugin

```
rabbitmq-plugins enable rabbitmq_federation
```



Enabling the Plugin

```
rabbitmq-plugins enable rabbitmq_federation
```

```
rabbitmq-plugins enable rabbitmq_federation_management
```



Federating an Exchange

```
rabbitmqctl set_parameter federation-upstream my-upstream \  
'{"uri": "amqp://server-name", "expires": 3600000}'
```



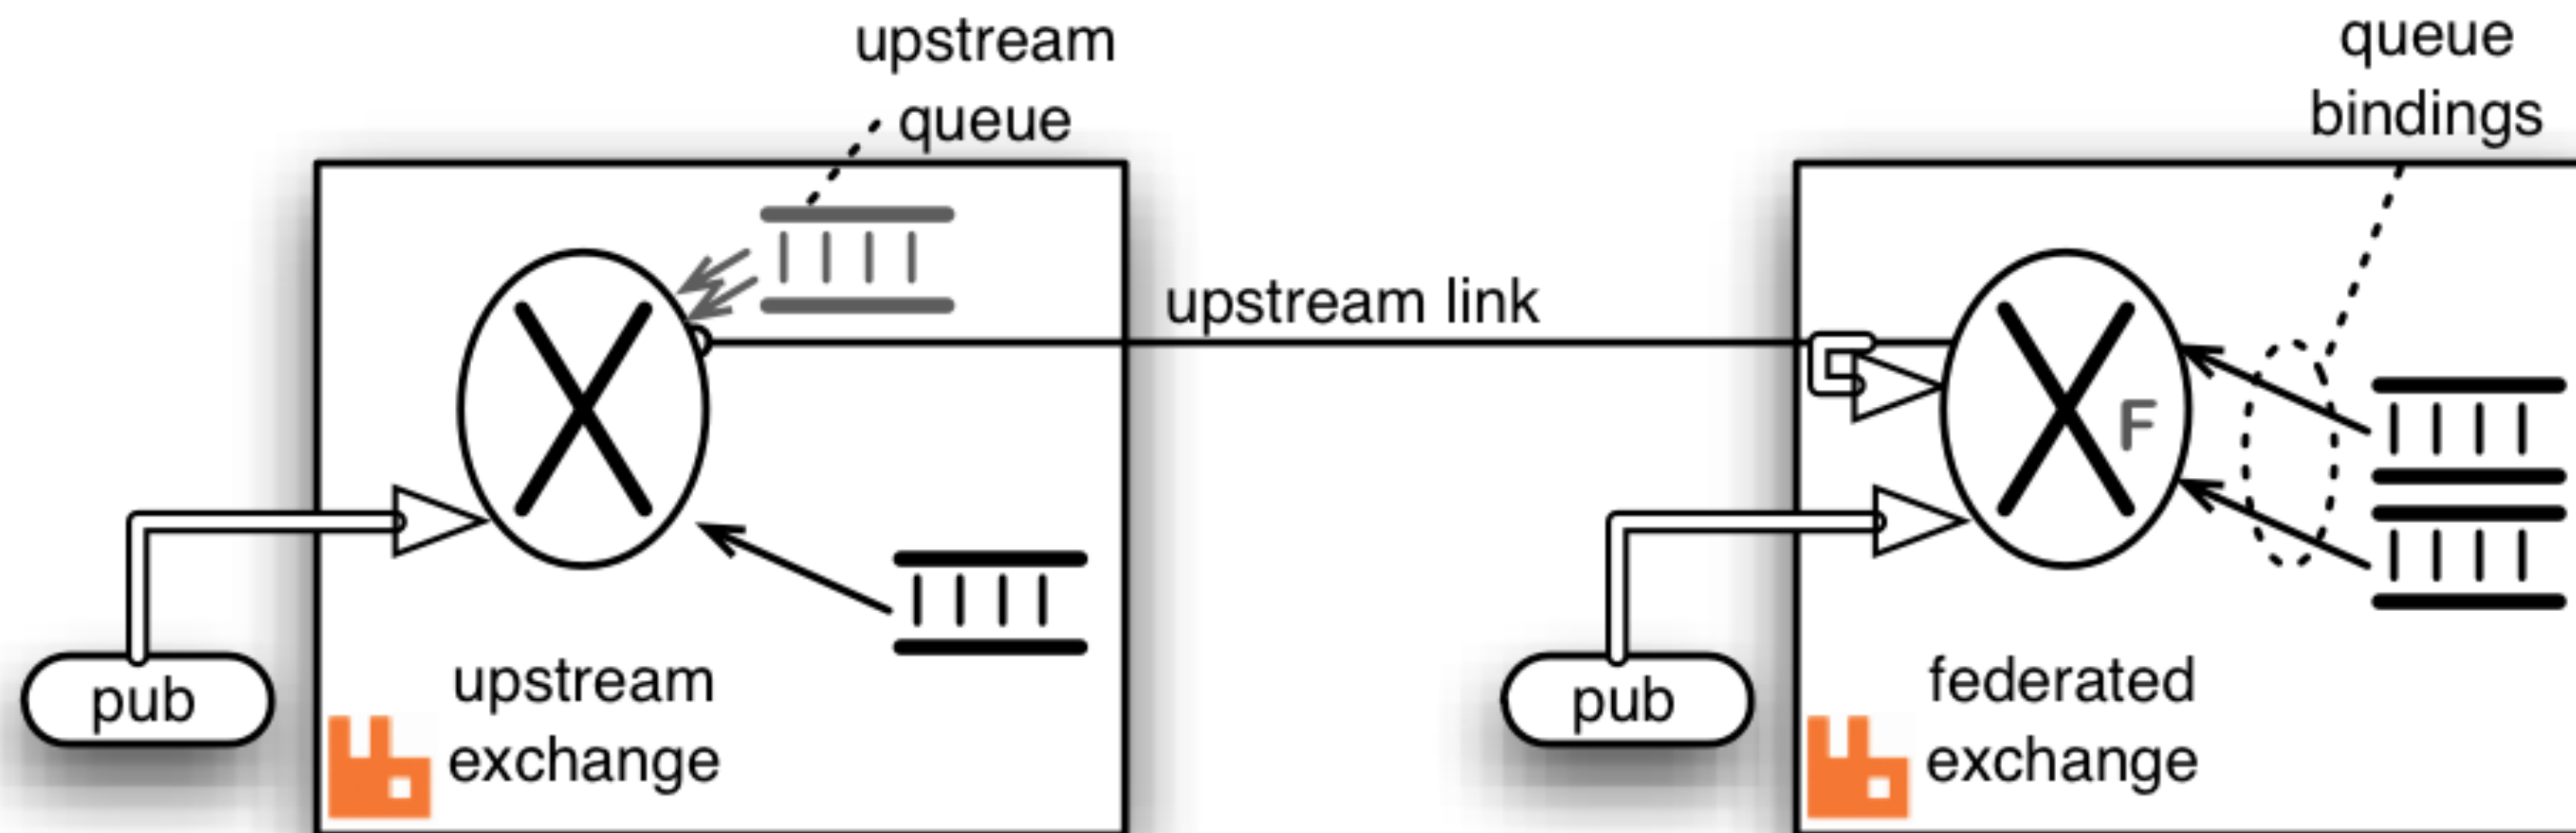
Federating an Exchange

```
rabbitmqctl set_parameter federation-upstream my-upstream \  
'{"uri": "amqp://server-name", "expires": 3600000}'
```

```
rabbitmqctl set_policy --apply-to exchanges federate-me "^amq\." \  
'{"federation-upstream-set": "all"}'
```



Federating an Exchange



Configuring Federation



Config Options

```
rabbitmqctl set_parameter federation-upstream \  
name 'json-object'
```



Config Options

```
rabbitmqctl set_parameter federation-upstream \  
name 'json-object'
```

```
json-object: {  
  'uri': 'amqp://server-name/',  
  'prefetch-count': 1000,  
  'reconnect-delay': 1,  
  'ack-mode': on-confirm  
}
```

<http://www.rabbitmq.com/federation-reference.html>



Prevent unbound buffers

`expires: N // ms.`

`message-ttl: N // ms.`



Prevent message forwarding

`max-hops: N`



Speed vs No Message Loss

ack-mode: on-confirm

ack-mode: on-publish

ack-mode: no-ack



AMQP URI:

`amqp://user:pass@host:10000/vhost`

<http://www.rabbitmq.com/uri-spec.html>

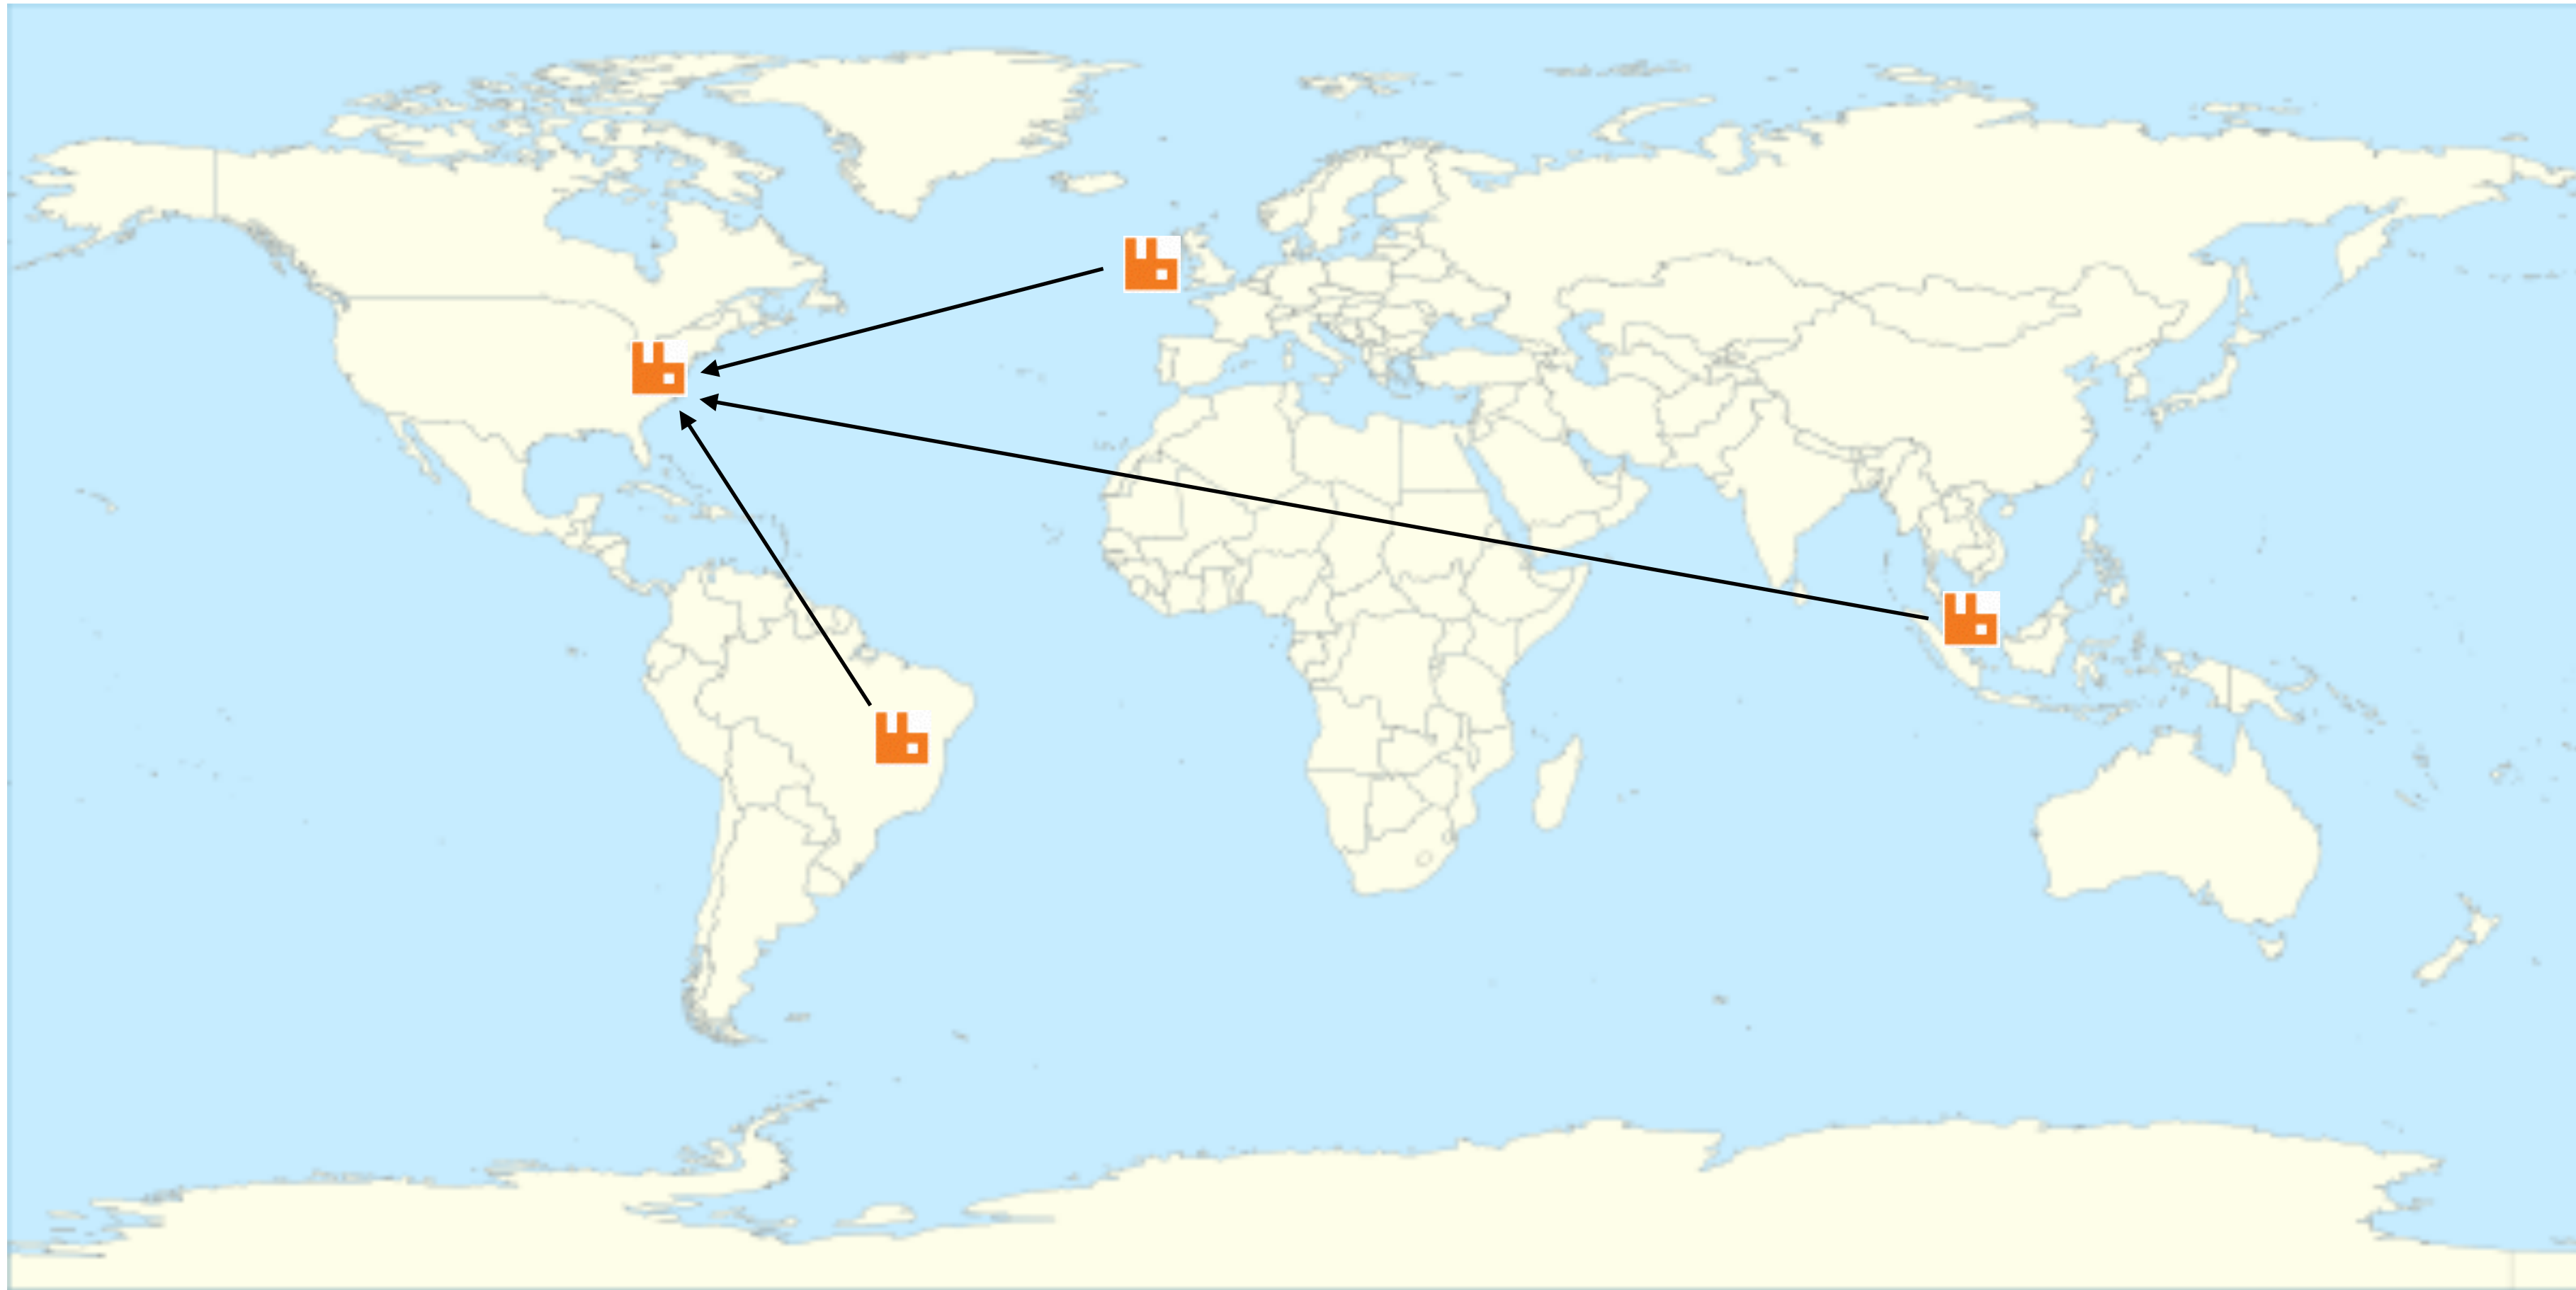


Config can be applied via

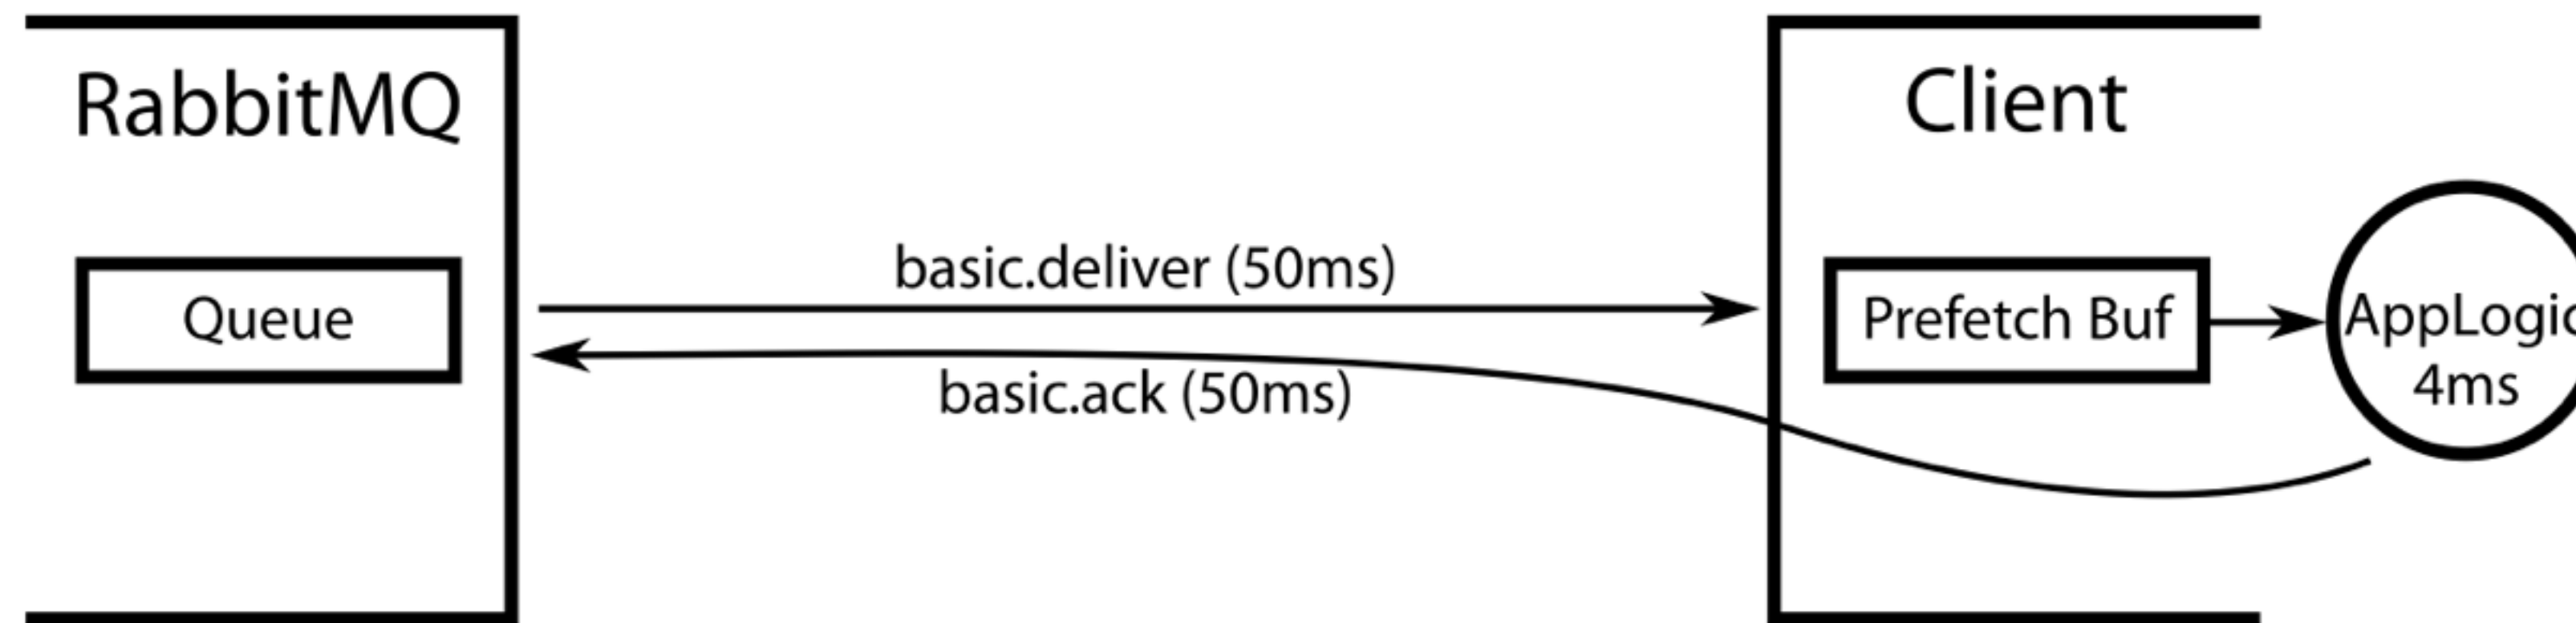
- CLI using **rabbitmqctl**
- HTTP API
- RabbitMQ Management Interface



RabbitMQ Federation

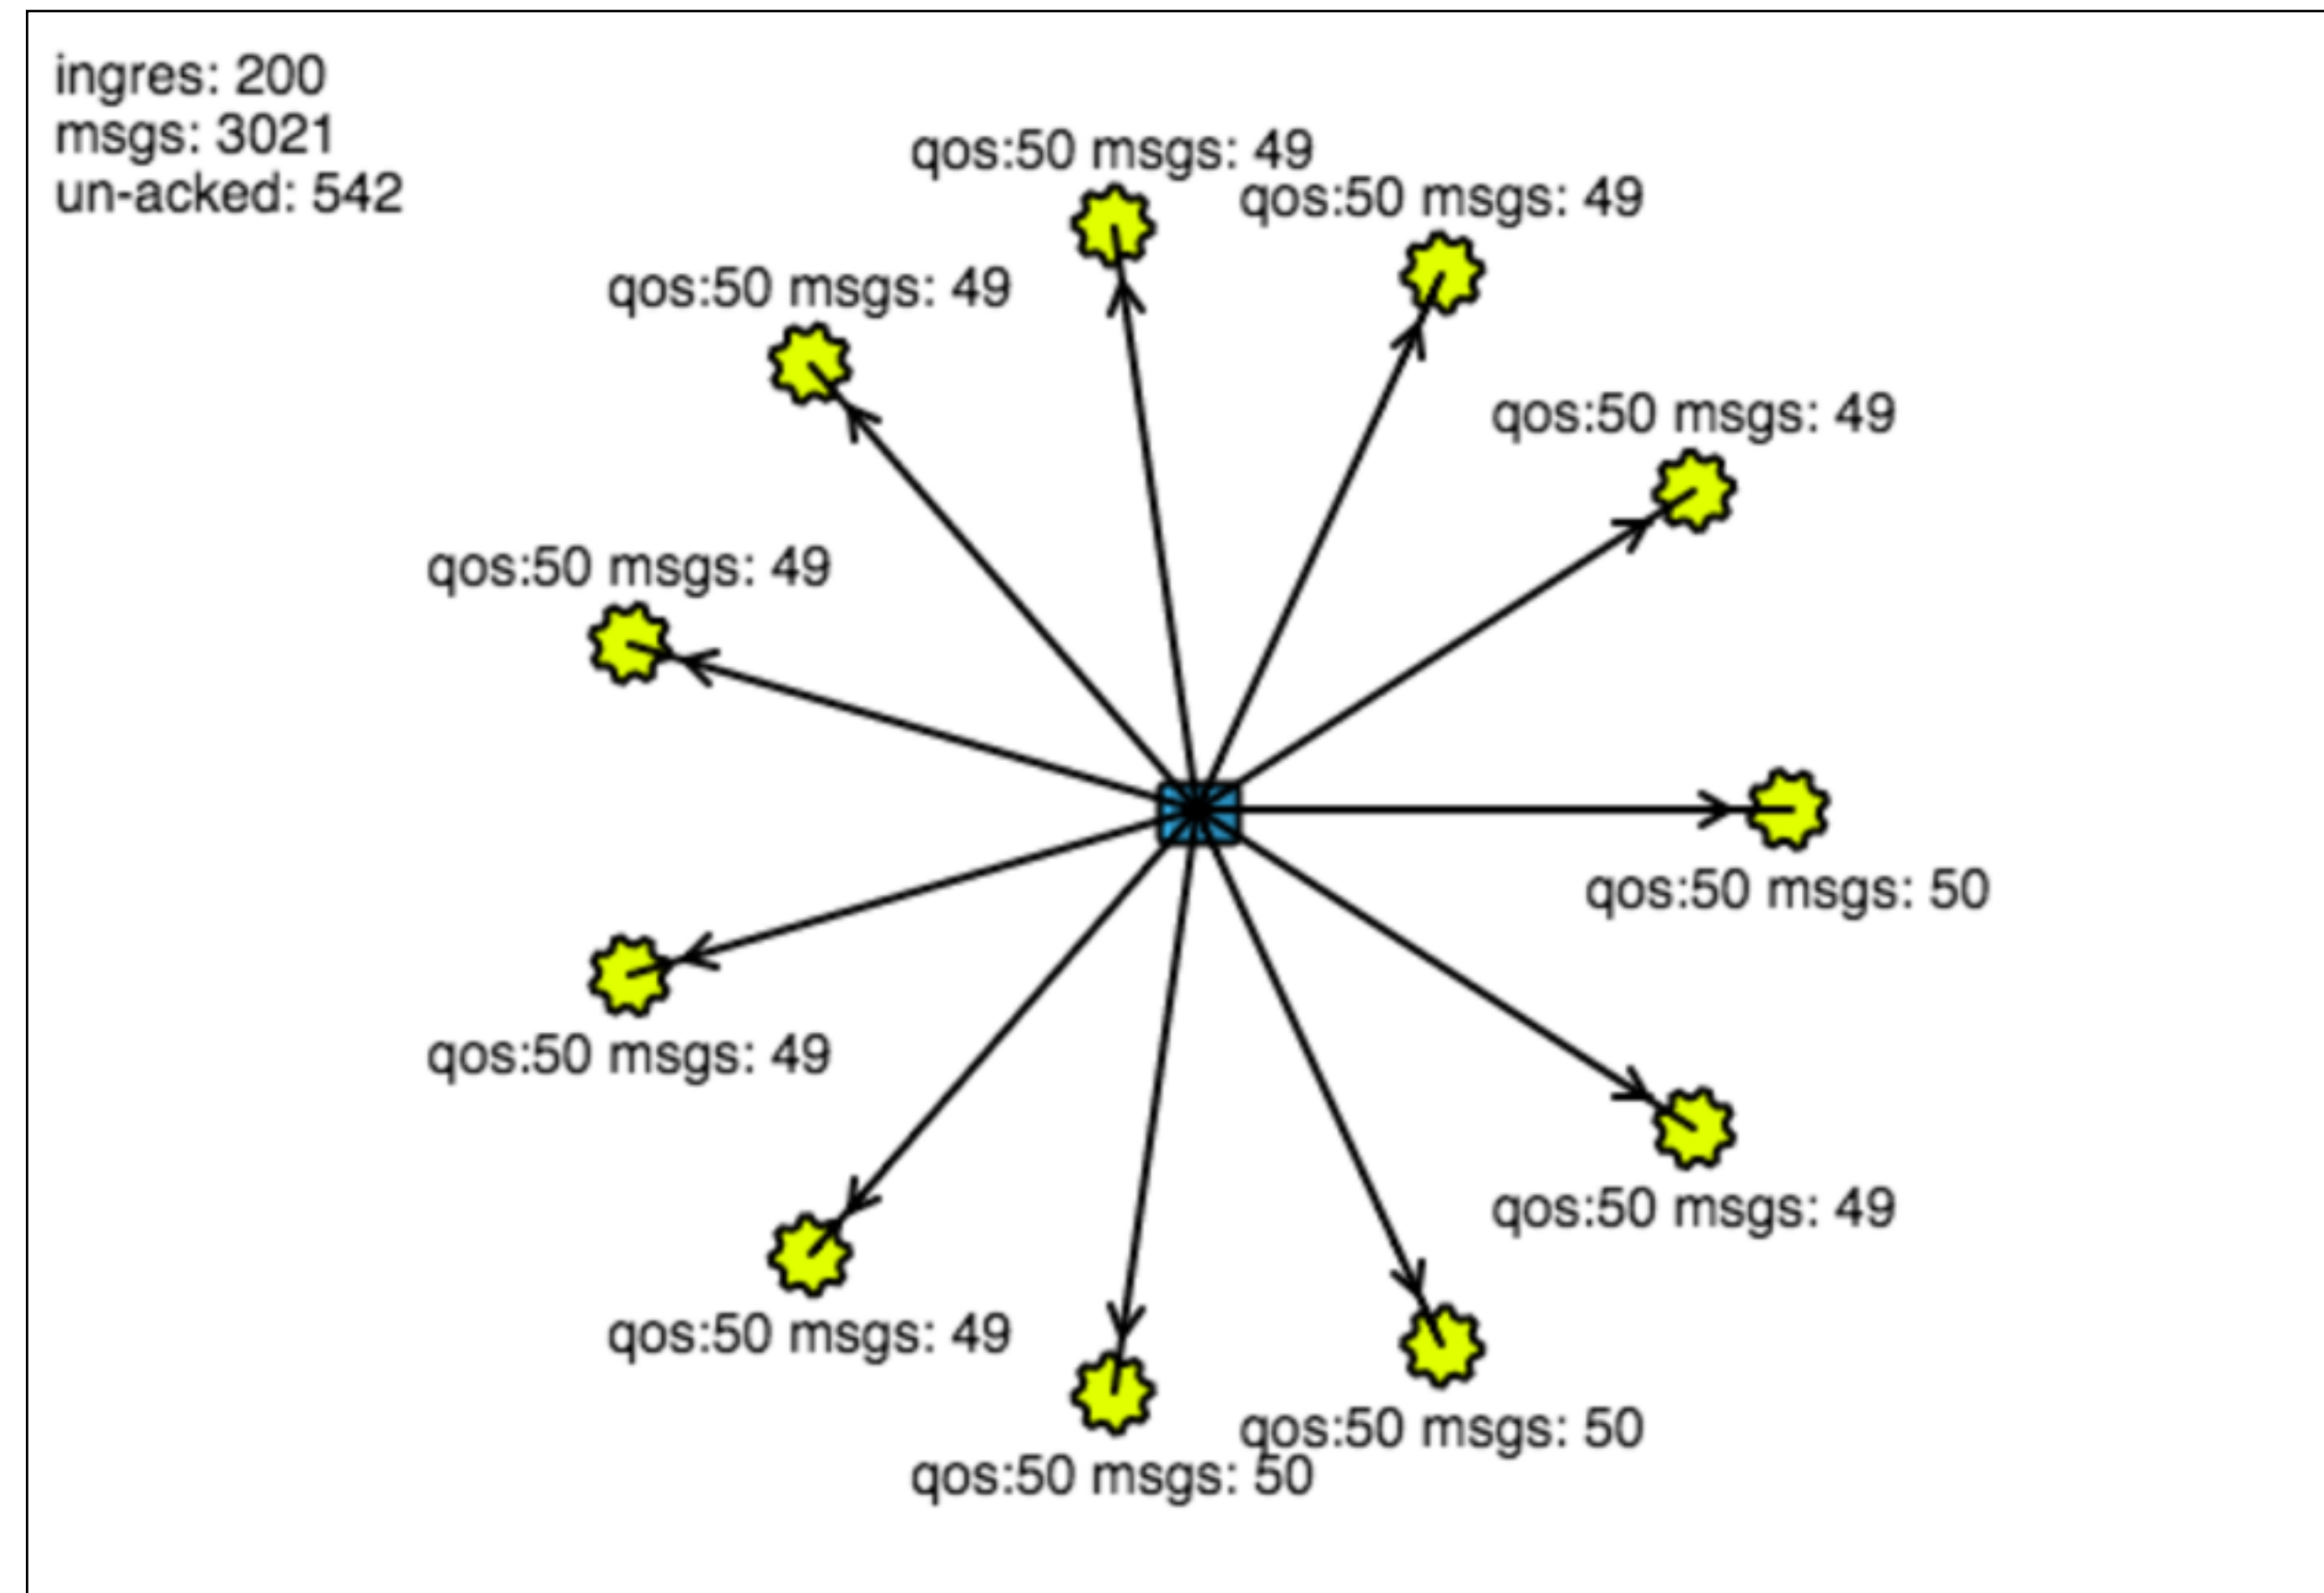


Some Queueing Theory



<http://www.rabbitmq.com/blog/2012/05/11/some-queueing-theory-throughput-latency-and-bandwidth/>

RabbitMQ BasicQos Simulator



Prevent Unbound Buffers

λ = mean arrival time

μ = mean service rate

if $\lambda > \mu$ what happens?

<https://www.rabbitmq.com/blog/2014/01/23/preventing-unbounded-buffers-with-rabbitmq/>



Prevent Unbound Buffers

λ = mean arrival time

μ = mean service rate


if $\lambda > \mu$ what happens?

Queue length goes to infinity over time.

<https://www.rabbitmq.com/blog/2014/01/23/preventing-unbounded-buffers-with-rabbitmq/>



Recommended Reading



**Performance Modeling and
Design of Computer Systems:
Queueing Theory in Action**

Scaling the Setup



The Problem



The Problem

- Queues contents live in the node where the Queue was declared



The Problem

- Queues contents live in the node where the Queue was declared
- A cluster can access the queue from every connected node



The Problem

- Queues contents live in the node where the Queue was declared
- A cluster can access the queue from every connected node
- Queues are an Erlang process (tied to one core)



The Problem

- Queues contents live in the node where the Queue was declared
- A cluster can access the queue from every connected node
- Queues are an Erlang process (tied to one core)
- Adding more nodes doesn't really help



Enter Sharded Queues

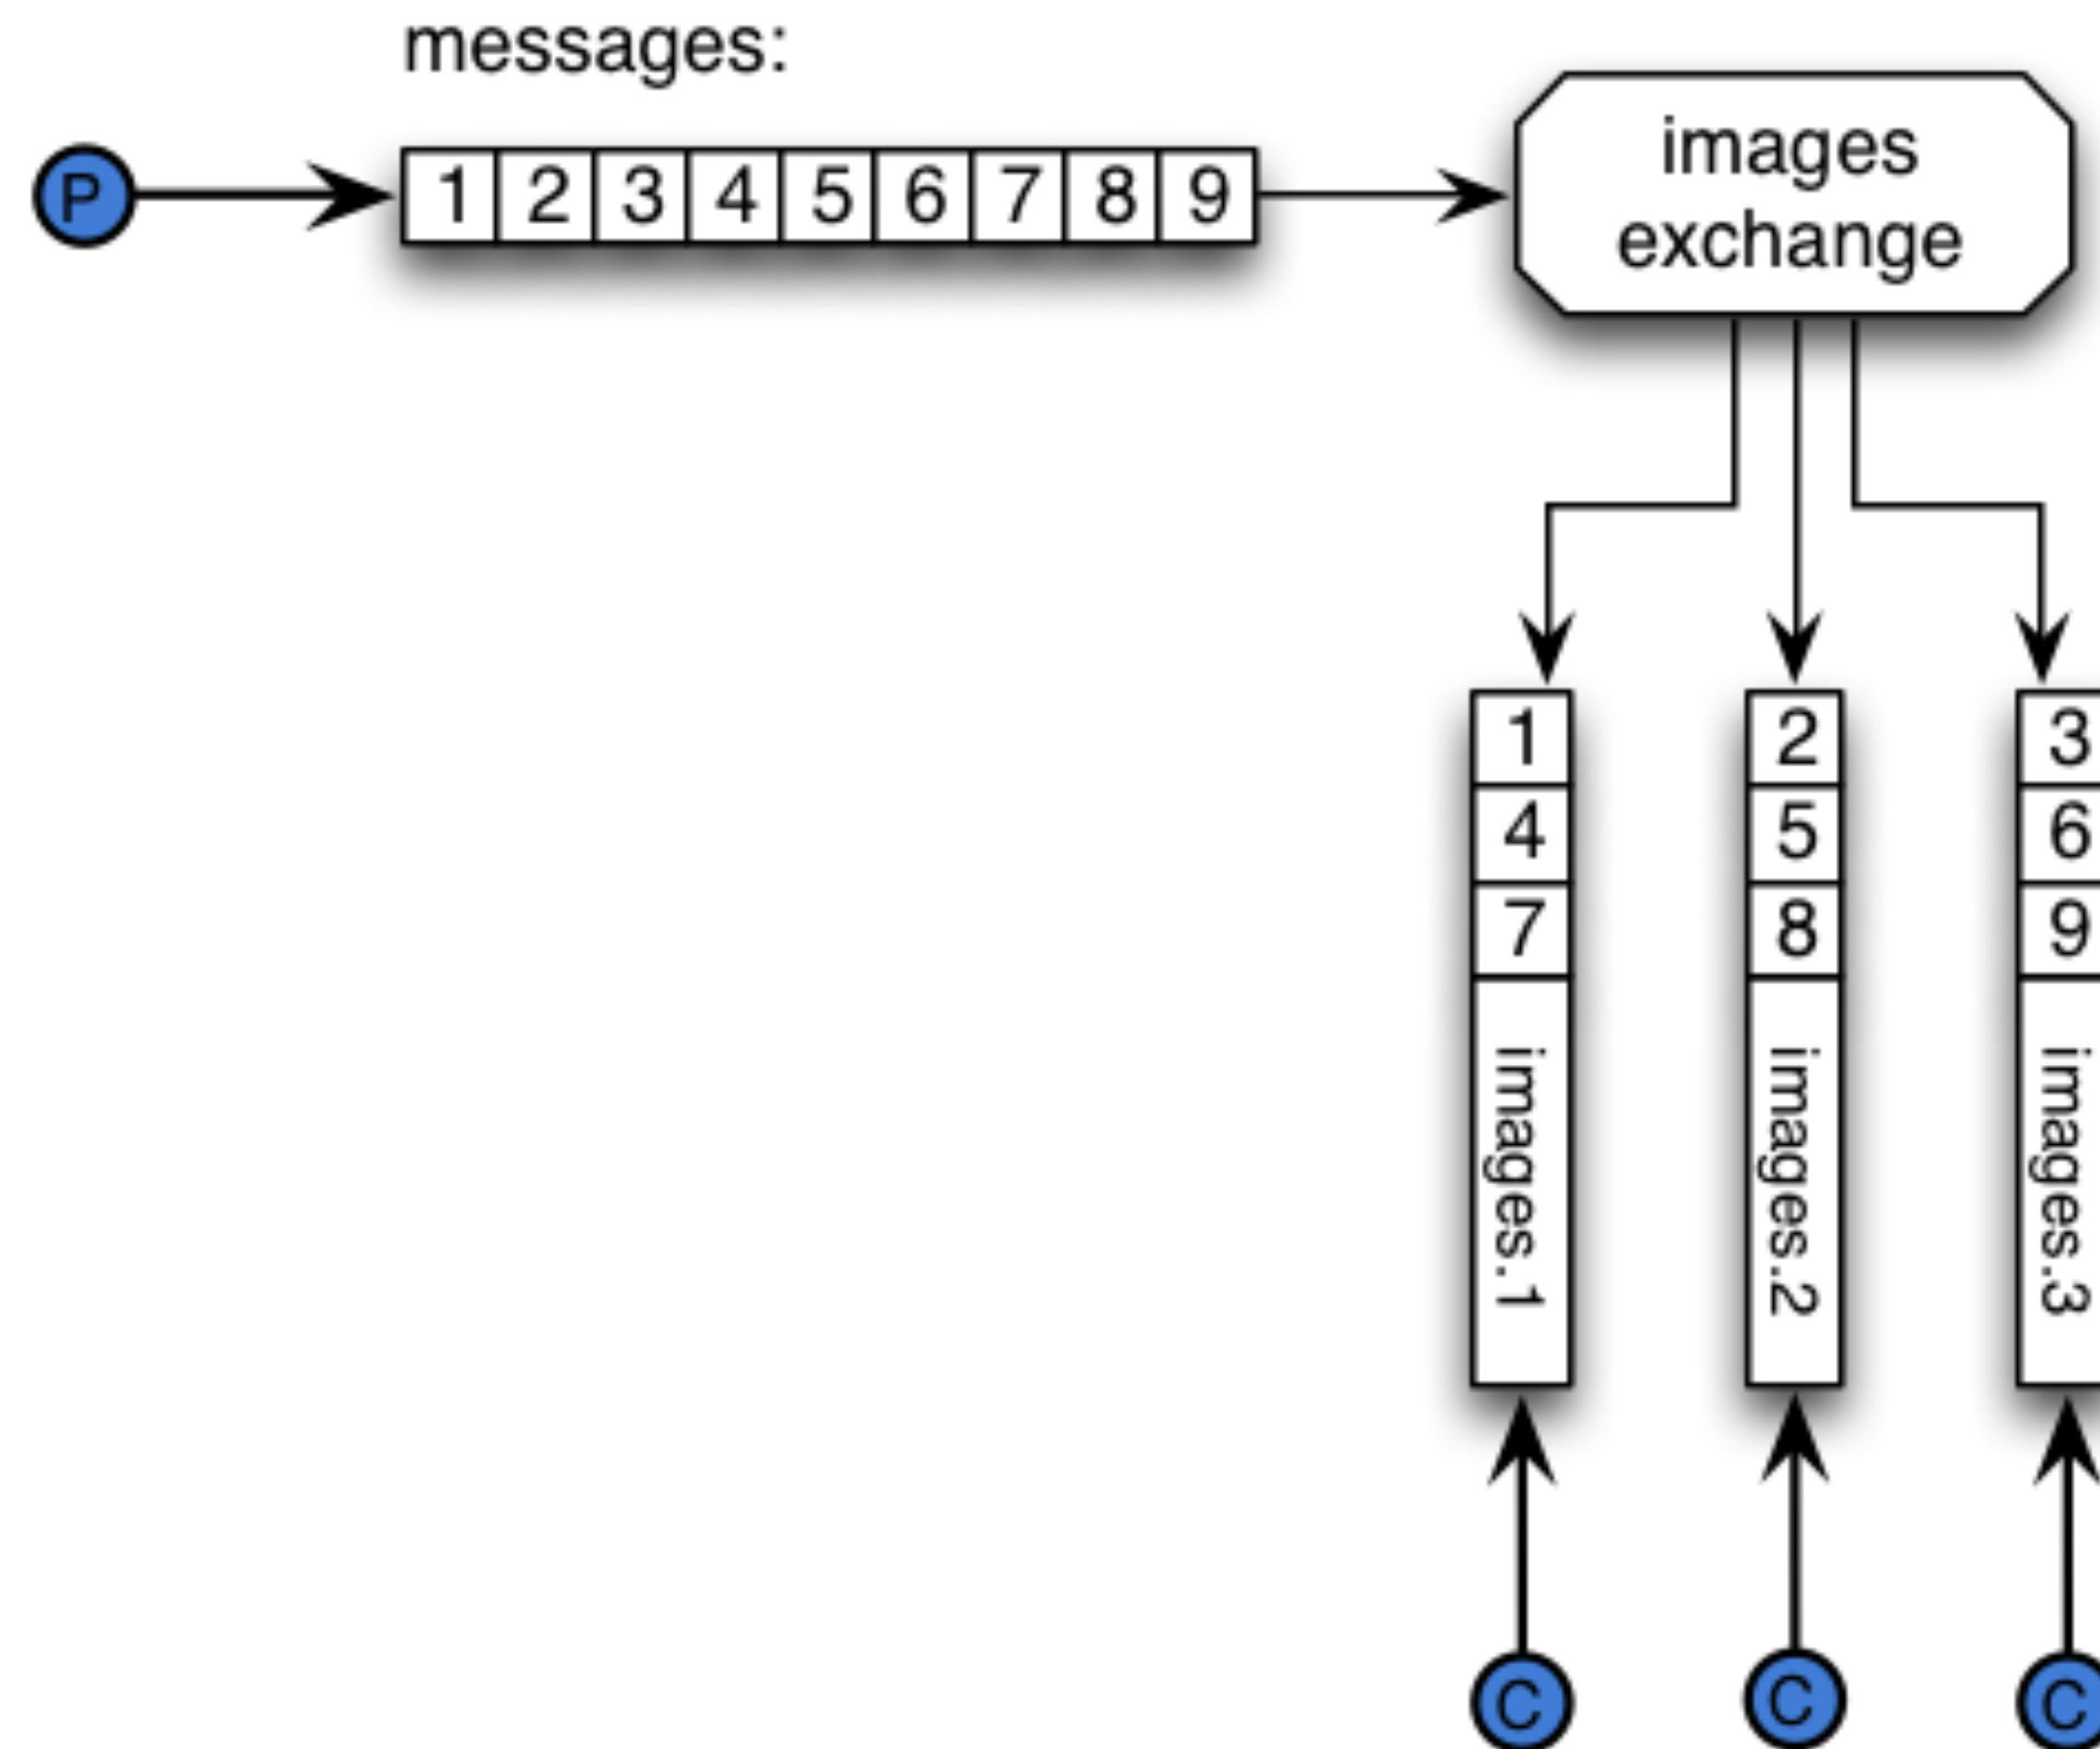


Pieces of the Puzzle

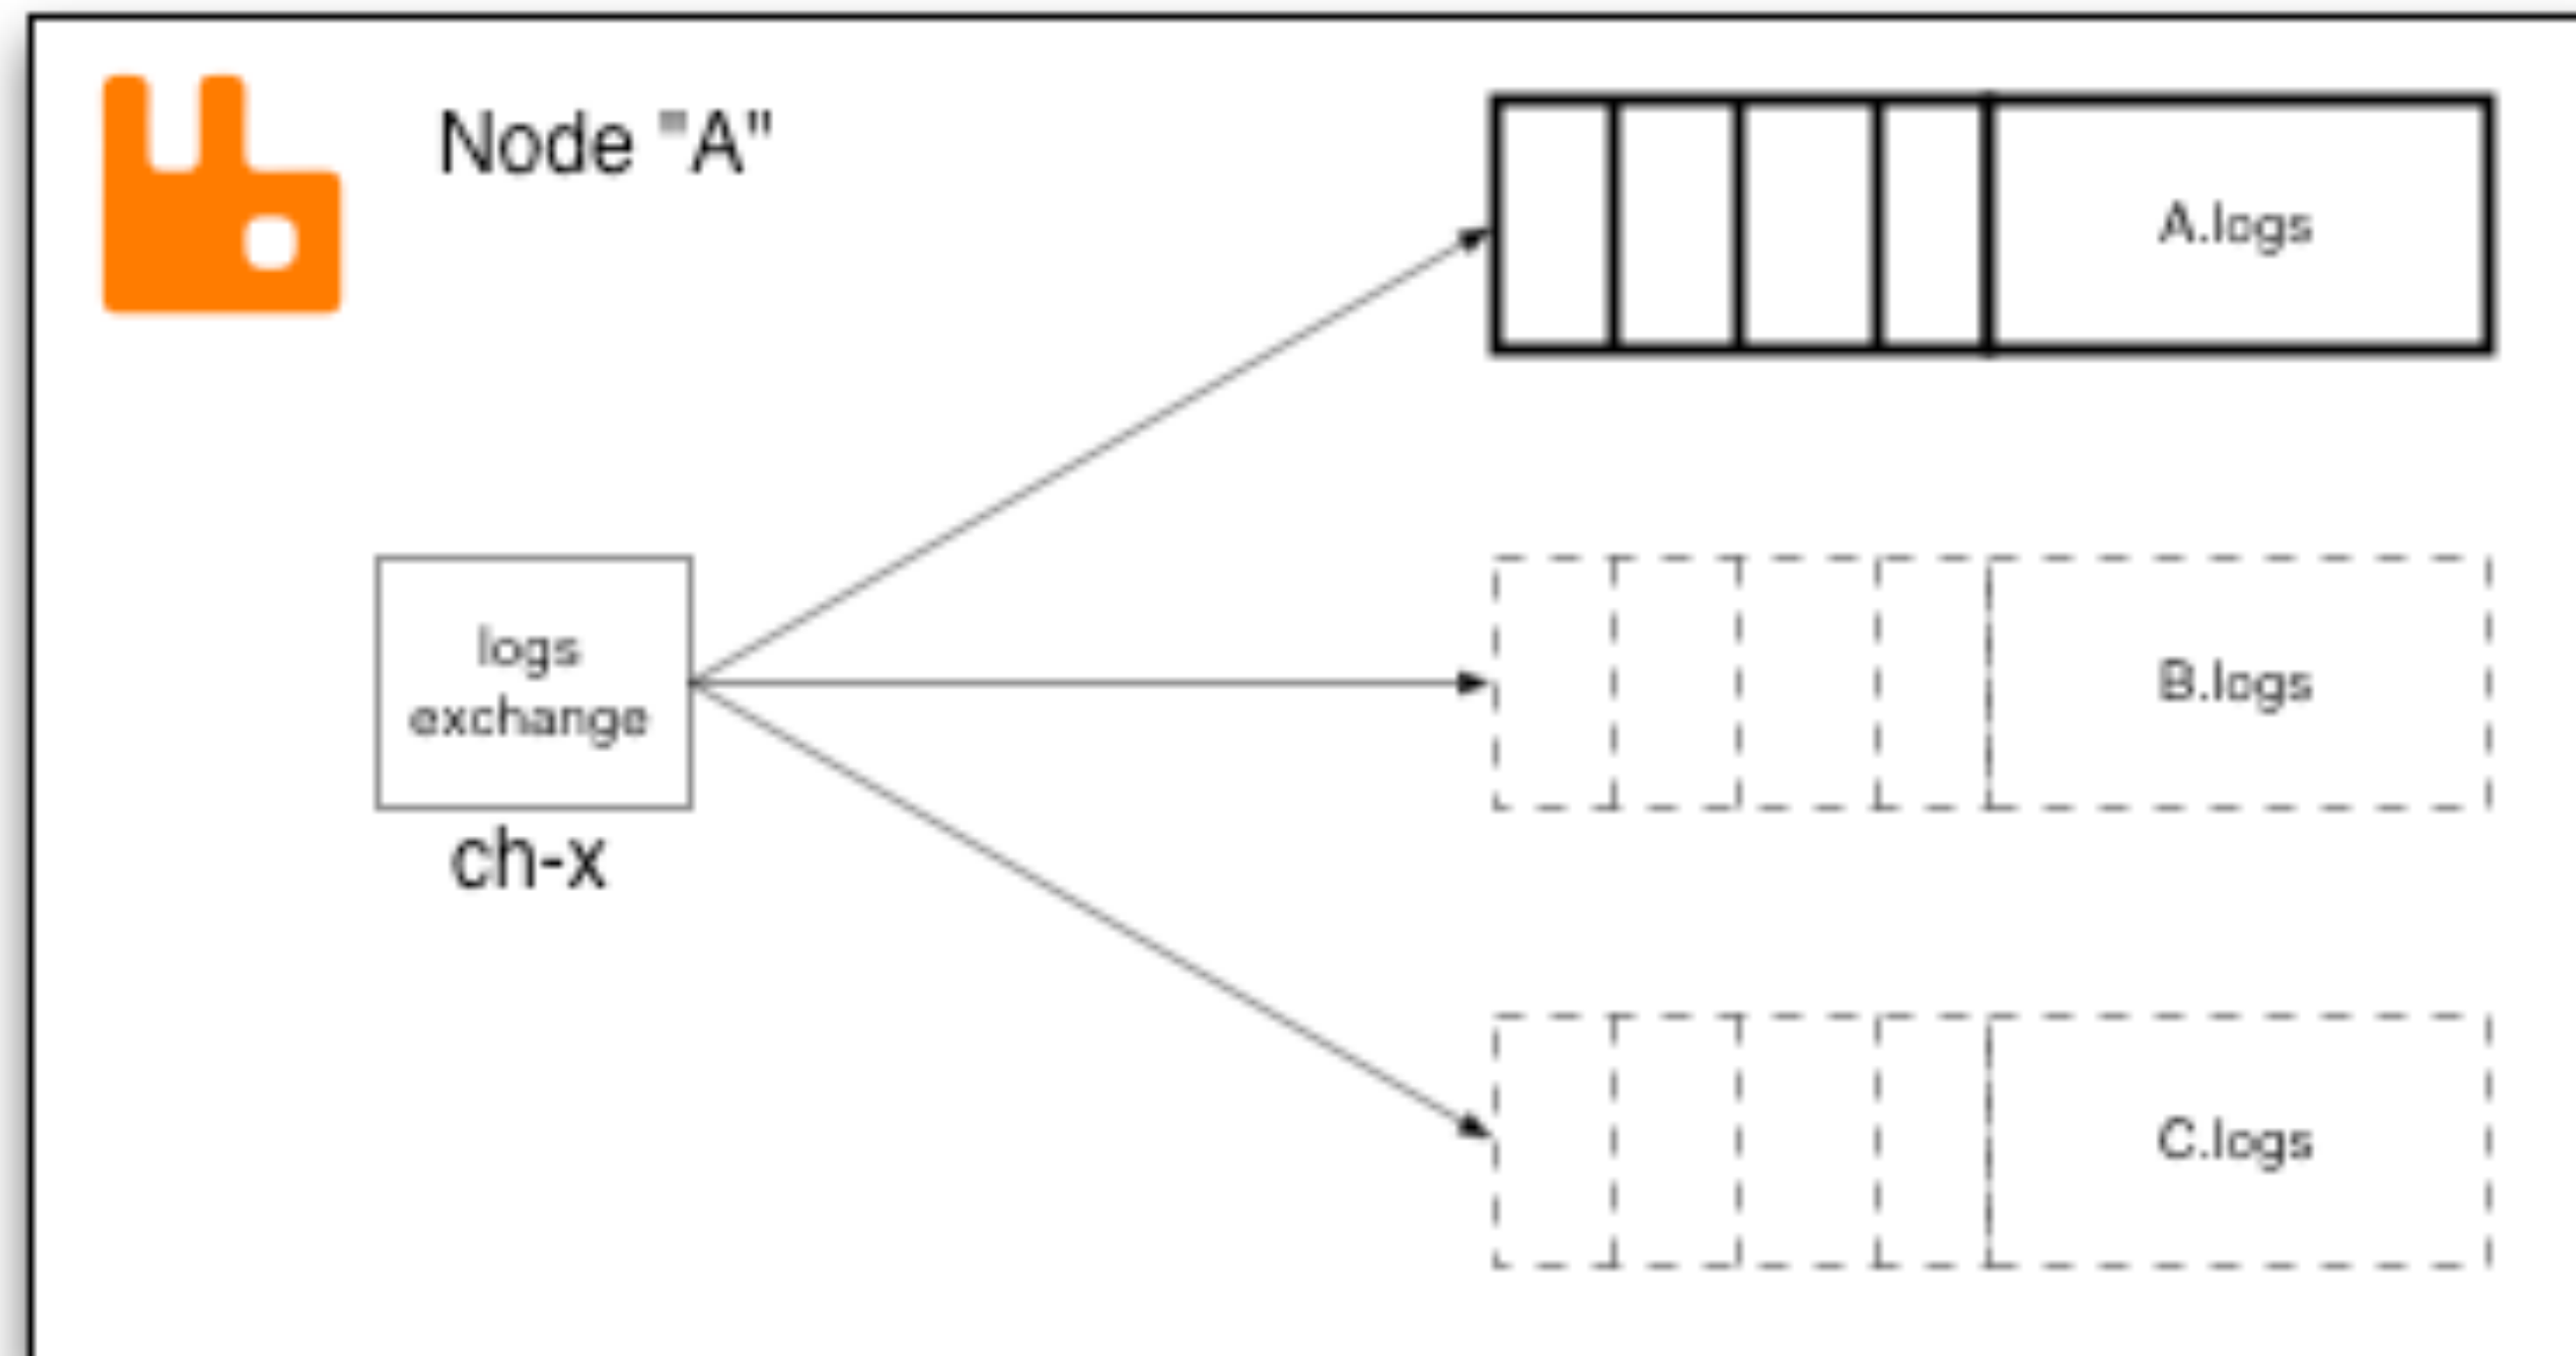
- modulo hash exchange (consistent hash works as well)
- good ol' queues




Sharded Queues




Sharded Queues



Sharded Queues



Sharded Queues



Sharded Queues

- Declare Queues with name: `nodename.queueName.index`



Sharded Queues

- Declare Queues with name: `nodename.queueName.index`
- Bind the queues to a partitioner exchange



Sharded Queues

- Declare Queues with name: `nodename.queueName.index`
- Bind the queues to a partitioner exchange
- Transparent to the consumer (virtual queue name)



We need more scale!



Federated Queues



Federated Queues

- Load-balance messages across federated queues



Federated Queues

- Load-balance messages across federated queues
- Only moves messages when needed



Federating a Queue

```
rabbitmqctl set_parameter federation-upstream my-upstream \  
'{"uri": "amqp://server-name", "expires": 3600000}'
```



Federating a Queue

```
rabbitmqctl set_parameter federation-upstream my-upstream \  
'{"uri": "amqp://server-name", "expires": 3600000}'
```

```
rabbitmqctl set_policy --apply-to queues federate-me "^images\." \  
'{"federation-upstream-set": "all"}'
```



With RabbitMQ we can



With RabbitMQ we can

- Ingest data using various protocols: AMQP, MQTT and STOMP



With RabbitMQ we can

- Ingest data using various protocols: AMQP, MQTT and STOMP
- Distribute that data globally using Federation



With RabbitMQ we can

- Ingest data using various protocols: AMQP, MQTT and STOMP
- Distribute that data globally using Federation
- Scale up using Sharding



With RabbitMQ we can

- Ingest data using various protocols: AMQP, MQTT and STOMP
- Distribute that data globally using Federation
- Scale up using Sharding
- Load balance consumers with Federated Queues



Credits

world map: wikipedia.org

federation diagrams: rabbitmq.com



Questions?



Конференция разработчиков
высоконагруженных систем



Thanks!

Alvaro Videla - @old_sound



Конференция разработчиков
высоконагруженных систем

